

class07

Mina Wu (A59013200)


2024-01-31

Today we are going to explore some core machine learning methods. Namely clustering and dimensionality reduction approaches.

Kmeans clustering

The main function for k-means in “base” R is called `kmeans()`. Let’s first make up some data to see how kmeans works and to get at the results.

```
hist(rnorm(50000, mean=3))
```




Make a vector with 60 total points half centered at +3 and half centered at -3.

```
tmp <- c(rnorm(30, mean= 3), rnorm(30, mean= -3))
tmp
```

```
[1] 2.9462402 1.6831047 2.1060236 2.0414446 4.4500104 2.9133869
[7] 1.8971422 2.9606544 1.8619032 3.4854480 1.0572305 2.9949316
[13] 3.8269230 4.1566482 0.9460158 3.2814959 5.1550833 4.2691873
[19] 2.8757692 2.6623884 4.2165154 2.2137103 1.5276517 2.4680147
[25] 4.2266682 3.4923890 2.1357236 3.0734172 2.5332365 3.0394980
[31] -4.3766072 -4.5340105 -3.8693593 -2.0692917 -3.3933618 -1.1001413
[37] -2.0221615 -4.1444268 -2.7951515 -2.8920546 -3.8284064 -2.6651104
[43] -3.1982104 -4.3640358 -2.5610630 -3.1346082 -4.5362506 -3.8248224
[49] -3.0666335 -2.0881991 -5.1340434 -3.7430728 -4.1304933 -2.6946165
[55] -5.6758845 -3.6552423 -2.7716627 -4.3657064 -4.3356361 -2.2133748
```

reverse the order of tmp using rev() for y. Then plot using plot().

```
x<- cbind(x=tmp, y=rev(tmp))
plot(x)
```



Run kmeans() asking for two clusters: and nstart=20, which is the iterations

K

C

12

C

C
IW
[

W
[

A

[]

V

\$11

\$1

V

12


What is my clustering results? I.E. what cluster does each point reside in?

```
k$cluster
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2  
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```


Q. Plot your data `x` showing your clustering result and the center point for each cluster?

```
plot(x, col= k$cluster)  
points(k$centers, pch=15, col = "green")
```



Q. Run kmeans and cluster into 3 groups and plot the result?

```
k3 <- kmeans(x, centers=3)  
plot(x, col=k3$cluster)
```



```
k$tot.withinss
```

```
[1] 125.9693
```

```
k3$tot.withinss
```

```
[1] 101.1119
```

The big limitation of kmeans is that it imposes a structure on our data (i.e. a clustering) that you ask for in the first place. K-means will always give you the renumber of clusters you request.

hierarchical clustering

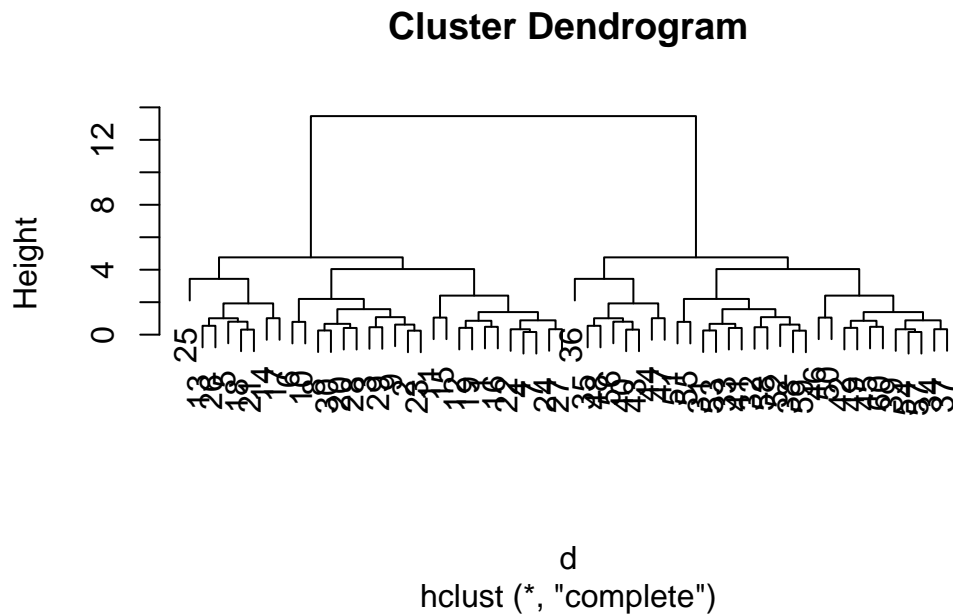
The main function in "base" R for this is called 'hclust()'. It wants a distance matrix as input no the data itself.

We can calculate a distance matrix in lots of different ways but here we will use the 'dist()' function.

```
d <- dist(x, diag = T)
hc<- hclust(d)
```

There is a specific plot method for hclust objects. Let's see it:


```
plot(hc)
```



cut the cluster at height 9 using abline.

```
plot(hc)
abline(h=9, col="red")
```

Cluster Dendrogram



d
hclust (*, "complete")

To get the cluster membership vector we need to “cut” the tree at a given height that we pick. The function to do this is called `cutree()`

```
cutree(hc, h=9)
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```


Use `k` to cut into 4 clusters.

```
grps<- cutree(hc, k=4)
grps
```

```
[1] 1 1 1 1 2 1 1 1 1 1 1 1 2 2 1 1 2 2 1 1 2 1 1 1 2 2 1 1 1 1 3 3 3 3 4 4 3 3
[39] 3 4 3 3 4 4 3 3 4 4 3 3 3 3 3 3 3 3 4 3 3 3 3 3 4 3 3 3 3 3 3 3 3 3 3 3 3 3
```

Q. Plot our data (`x`) colored by our `hclust` result.

```
plot(x, col = grps)
```



Principal Component Analysis (PCA)

We will start with PCR of a tiny tiny dataset and make fun of stuff Barry eats.

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names=1)
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494

Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41


```
dim(x)
```

```
[1] 17  4
```

```
# there are 17 rows and 4 columns.
```

One useful plot in this case (because we only have 4 countries to look across) is a so-called pairs plot.

```
pairs(x, col=rainbow(10), pch=16)
```



Enter PCA

The main function to do PCA in “base” R is called `prcomp()`.

It wants our foods as the columns and the countries as the rows. It basically wants the transpose of the data we have.

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	4.189e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
attributes(pca)
```

\$names

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

\$class

```
[1] "prcomp"
```

Plot the PCA plot

```
plot(pca$x[,1], pca$x[,2], xlab= "PC1(67.4%)", ylab= "PC2(29%)", col= c("orange", "red", "gray"),
     abline(h=0, col="gray", lty=2)
     abline(v=0, col="gray", lty=2))
```

